

# High Performance LLMs From First Principles (2024)

**Goal: learn how to achieve high  
performance for LLMs**

**This week:**

- **End-to-end Training, Measure Performance**

**Program** (write code in Jax)

**Predict** (roofline on napkin or spreadsheet)

**Profile** (run code, compare to predictions)

# My Asks

Please ask lots of questions! Just raise your hand or speak up!

If there are topics you're interested in, message me between sessions.

Join the discord! <https://discord.gg/2AWcVatVAw>

Do the exercises! Give feedback, ask questions!

Website: <https://github.com/rwitten/HighPerfLLMs2024>

# Add ons to Session 1

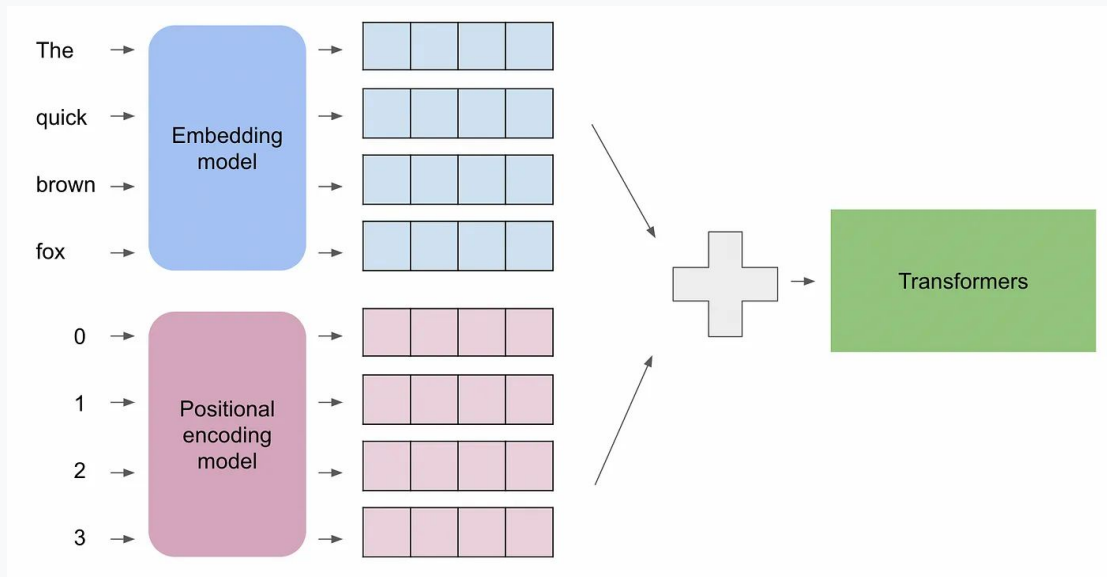
- Didn't have Attention
- Didn't do positional encoding
- Didn't JIT
- Overlapping Data Loading
- Was single chip (not multichip)
- Never computed model flops per second since we didn't yet know how

# Add Attention

- Code

# Add Positional Embedding

- The problem is that attention was order invariant
- “Brown” sees “the” and “quick” but doesn’t know the order!
- Solution – add embedding vector





# JIT

- Not just JITing forwards pass – want to JIT a whole computation traditionally!

# Overlapping Data Loading (Lay Track Faster Than Train!)

Naive:




Overlapped:



- Easy to do it wrong by accident! (Also easy to get it right!)
- Assuming you do it right:
  - Data loading doesn't matter unless it becomes the bottleneck
  - Then it matters a great deal because we're in 1 for 1 slip.

Overlapped slow:




# Multichip

- Need to generate params on multiple chips, generate data for multiple chips and then JIT it all together.
- Lots of options but actually a little trickier than it sounds:
  - The key issue is making sure to describe a distributed computation.
  - The anti-pattern is putting everything on one chip and then moving it around – this out-of-memories when things don't fit on a single chip!
  - Useful: `jax.ShapeDtypeStruct`

# How long should training take?

- We discussed the forwards pass in the past.
- Training adds the backwards pass! Backwards pass is computing the derivatives!


# Forwards Pass




# Corresponding Backwards Pass

Proprietary + Confidential

**"Spine"**



**"Weight Gradients"**



# How long should training take?

- We discussed the forwards pass in the past.
- Training adds the backwards pass! Backwards pass is computing the derivatives!
- For every matmul in the forwards pass, two matmuls in the backwards pass:
  - “Spine” – generate the new partial derivations
  - “Weight gradients” – output the gradients we need!
  - All the matmuls take  $2*B*E_1*E_2$  flops since they all have a B dimension, an  $E_1$  and an  $E_2$  dimensions! So  $6*B*E_1*E_2$  total flops.
- Conveniently:
  - Parameters =  $E_1*E_2$
  - So this matmul during training takes  $6*B*P$  total flops.
  - And that is true for all matmuls – so the number of flops from matmuls is  $6*B*P$ !
- Warning: this calculation ignores attention flops! Attention flops are typically a small fraction of the overall flops...

# Next Week

- Get the model converging and do some inference.
- Lots of possible bonus topics, not sure what has interest, happy to cover a couple more if there is strong interest.
  - “Going a Level Deeper in shard\_map and Pallas” (possibly with Sharad!)



**Thanks!**

**Ping me (rwitten@google.com) with  
feedback, suggested topics, etc!**